

Prototype Scope

Overall Goal

- Demonstrate an end-to-end flow: **Image upload → AI assistive suggestion → structured output saved for analysis**
 - Prove technical feasibility of AI-assisted wetland observation without requiring user expertise
 - Produce a reviewable artifact that shows **valid system design for MVP - not full feature completeness**
-

Interpreting the 3 Guidelines

1) Backend only ✓ **The backend handles file input, API call, response parsing.**

- Focus on API logic, data handling, and AI interaction; **UI can be minimal or simulated**
- Prioritize correctness, clarity, and logging over polish
- Enables fast iteration and avoids frontend bottlenecks

2) Picture file upload in VS Code ✓ **You read a local image file and send it via the API.**

- Local file upload simulates future web uploads without building auth or storage
- Keeps scope contained: single image, predictable format, **repeatable tests**
- Allows deterministic debugging of AI responses

3) AI response with ID

- AI returns *suggested biological labels*, not verified classifications
 - Output must be structured (JSON), not free text
 - Response is explicitly assistive, not authoritative
-

AI Identification Design

AI should suggest 2-3 species/genus members

Why:

- Single ID implies false certainty (especially with non-expert users)
- 2-3 options reflect how real ecological review works
- Supports your *human-in-the-loop* and *responsible AI* framing

Best practice for a prototype:

- Primary suggestion (most likely)
- 1-2 alternative possibilities
- Each with a confidence estimate or qualitative ranking

Even better (if time allows):

- Species *if confident*, otherwise Genus-level
- Explicit fallback: “*Unable to confidently identify. Do you want ideas for better photos?*”

Using LLaMA Vision – Where the API Fits

You are not training a model. You are calling one.

Here's the clean mental model.

Architecture (Prototype-Appropriate)

[Local Image File]



[Backend Script / Service]



[LLaMA Vision API Call]



[Structured AI Response (JSON)]



[Console Output / Saved Record]

API's Role

→ Transmits the image + instructions and returns structured suggestions; no storage, no training, no state.

- The API is how your backend sends:
 - The image
 - A carefully written prompt (instructions + constraints)
- And receives:
 - The AI's structured response (IDs + confidence)

You are using the API as:

“A remote, pre-trained perception service.”

Example: What the AI Response Should Look Like (Conceptually)

```
{  
  "suggestions": [  
    {  
      "label": "Red-eared slider (Trachemys scripta elegans)",  
      "rank": 1,  
      "confidence": 0.62  
    },  
    {  
      "label": "Painted turtle (Chrysemys picta)",  
      "rank": 2,  
      "confidence": 0.24  
    }  
],  
  "disclaimer": "Suggestions only. Not a verified identification."  
}
```

Coach's Advice (Straight Talk)

With 8 days:

- Do not chase UI
- Do not chase accuracy perfection
- Do not oversell automation

What *will* impress:

- Clean flow
 - Clear AI boundaries
 - Honest uncertainty
 - Structured outputs
-

If you want, next I can:

- Draft the exact LLaMA Vision prompt
- Sketch the backend pseudocode
- Help you write the “Prototype Limitations” section (very powerful for reviewers)